EBOOK

The Connected Enterprise AI Playbook

Moving from productivity tools to collaborative intelligence

The companies winning with AI aren't deploying smarter tools. They're building systems where humans and AI collaborate as a unified whole.

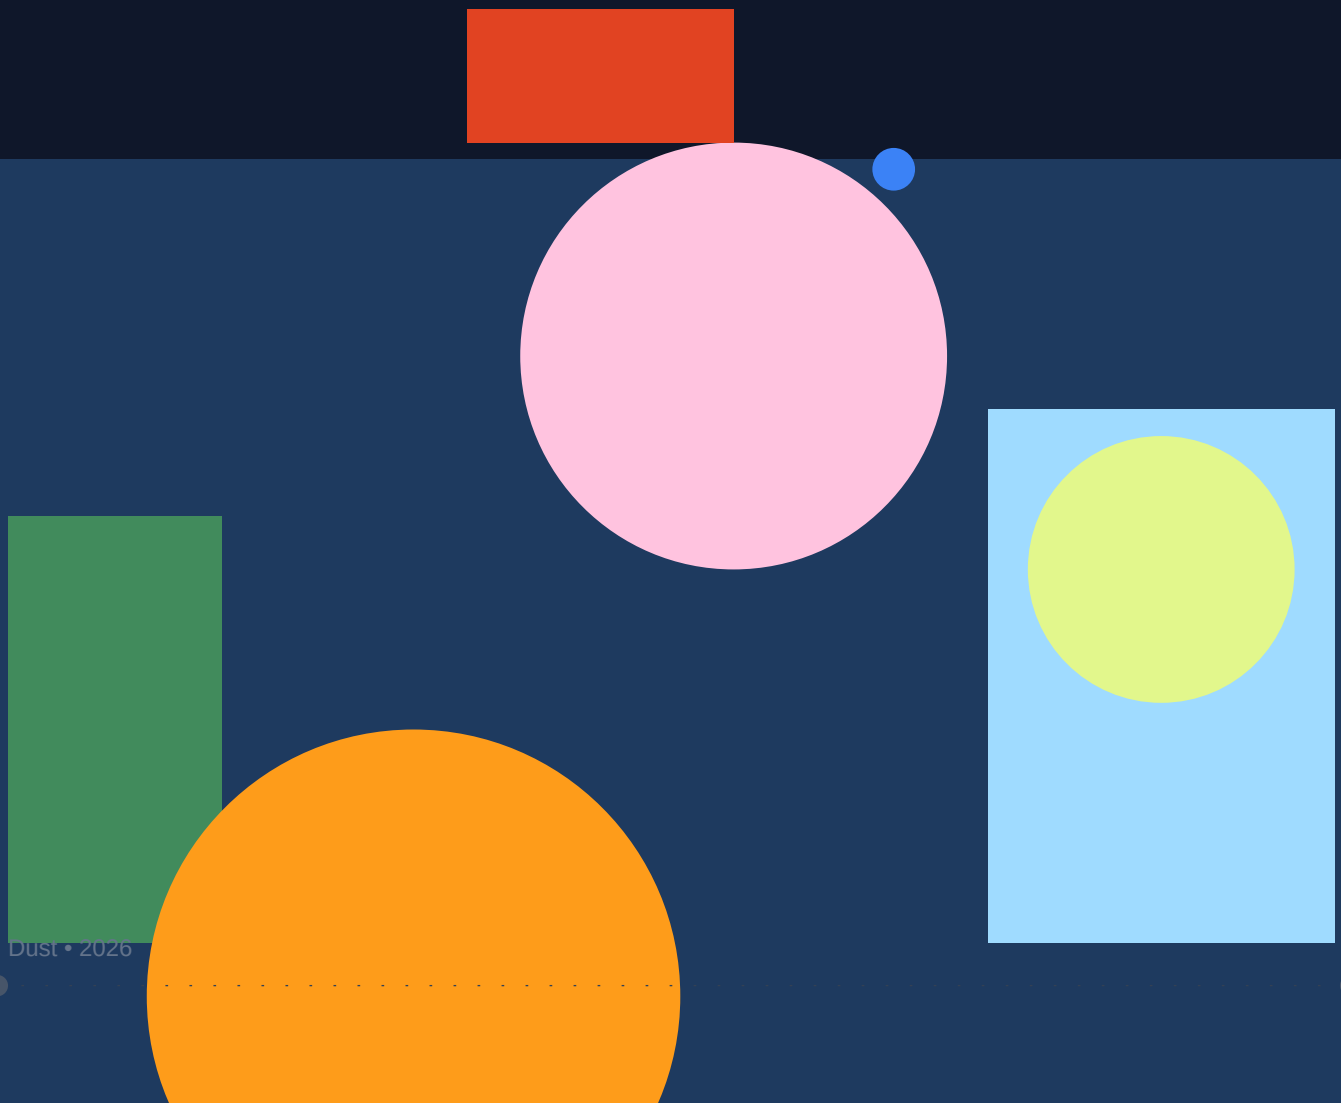

Dust • 2026

Table of contents

1 The shift: From individual AI to collaborative intelligence

2 The three waves of enterprise AI

3 What companies mistake for success

4 Core capabilities of collaborative intelligence

5 Implementation framework

6 Measuring success

7 The path forward

INTRODUCTION

The shift

> *"In 2024, your company bought AI to make employees more productive. In 2026, the question isn't whether AI makes individuals faster. It's whether it changes how your organization works. There's a difference. And that difference is where competitive advantage lives."*

Most enterprise AI conversations fixate on individual productivity gains. The real transformation is different: it's about how work moves between people, between teams, between humans and AI.

Companies deploying AI hit a ceiling. Adoption happens, productivity bumps occur, but organizational capability doesn't fundamentally change. Why? They're stuck in single-player mode when the game has gone multiplayer.

Single-player vs. multiplayer AI

Single-player AI: Each person has their own AI assistant. Insights stay isolated. Improvements don't spread. Knowledge fragments across tools and teams.

Multiplayer AI: Humans and AI collaborate together. Agents work with other agents. Knowledge compounds. The organization learns as a system.

What this playbook covers: A new framework for what enterprise AI actually means — not better tools for individuals, but collaborative intelligence infrastructure that makes your entire organization more capable.

Who this is for

You're an operations leader, AI implementation champion, or technical decision-maker who has moved past *"should we use AI?"* to *"how do we make AI transform how we work?"*

You've probably:

- Deployed ChatGPT Enterprise, Microsoft Copilot, Glean, or Gemini — saw adoption and productivity gains
- Connected enterprise search tools to your knowledge base — useful, but incomplete
- Built custom agents or workflows — working in pockets, not spreading
- Hit questions like: *"How do we scale beyond early adopters?"* or *"Why isn't this creating the transformation we expected?"*

Most enterprise AI content assumes you're either a Fortune 500 with armies of consultants and 18-month transformation roadmaps, or an individual looking for personal productivity hacks. You're neither.

You're at a company (500 to 10,000 employees) that needs enterprise-grade AI without enterprise overhead. You need AI that doesn't just make individuals faster but rewires how teams collaborate and how organizational knowledge compounds.

> **Key insight:** The companies pulling away aren't using better AI. They're using AI that collaborates differently — with humans, with other AI, across teams — in ways that make organizational intelligence compound over time. This is about architecture, organization, and communication. Not just algorithms.

SECTION 1

The three waves of enterprise AI

Wave 1: Enterprise search (2015–2023)

The first wave was about finding things. "Google for your company." End the "where did someone document that?" problem. Make tribal knowledge searchable.

It solved real problems: reduced time hunting for information, connected disparate data sources (Slack, Notion, Drive), and made onboarding faster.

> *"We made knowledge more findable. We didn't make it more usable."*

The ceiling: AI that finds answers but can't act on them. Read-only by design. Each search is isolated — no memory, no learning, no compounding.

Representative players: Glean, Guru, early Notion AI

Wave 2: Conversational AI and LLM assistants (2022–2025)

Wave 2 gave AI a voice, but not hands. Chat with AI like a colleague. Draft documents, summarize meetings, translate content. A personal productivity multiplier.

But this is where most companies are stuck. Three traps define the ceiling:

1. The verticalization trap. Teams deploy different AI tools — Sales uses Tool A, Marketing uses Tool B, Support uses Tool C. Each tool optimized for one function, none connected to others. Knowledge doesn't compound. Vendor lock-in creates silos. If AI becomes how we collaborate (like email or Slack), fragmenting it across verticalized tools undermines the whole point.

2. The crowd problem. Everyone has their own AI assistant, but assistants don't talk to each other. Sales AI doesn't know what Customer Success learned. Marketing insights don't flow to Product. Each conversation isolated, improvements don't spread.

3. The entropy trap. Every team builds custom agents, prompts, workflows. Knowledge fragments across dozens of custom GPTs or hundreds of Copilot agents. The more you build, the more fragmented it becomes. Marginal complexity keeps growing.

> **Key insight:** Wave 2 makes individuals smarter but doesn't make the organization smarter. It's single-player AI in a multiplayer world.

Representative players: ChatGPT Enterprise, Microsoft Copilot, Claude for Enterprise, Gemini for Workspace

Wave 3: Collaborative intelligence systems (2025+)

This isn't "better search" or "smarter chat." It's fundamentally different architecture.

The best human teams don't just work in parallel. They coordinate through shared objectives, awareness of what teammates are doing, and moving in cohesion. Wave 3 brings that dynamic to how humans and AI collaborate.

Three defining characteristics:

1. Agent ↔ agent ↔ human ↔ human collaboration

AI agents collaborate with each other, with individual humans, and with multiple humans. Work flows between agents like it flows between team members. Agents can ping humans just as humans ping agents.

In practice: specialized agents you @mention like colleagues; sub-agents that spawn other agents for multi-step workflows; skills as shared capabilities any agent can invoke; projects as hubs where humans, agents, context, and goals converge.

2. Write actions, not just read

AI that can execute: update CRM, send emails, create documents, trigger workflows. Permission-aware, respecting access controls. Closes the loop: find insight → take action → learn from outcome.

Skills give agents more "brain" (knowledge and reasoning). Tools (MCPs) give agents "hands" to perform actions across your systems. Together: agents that can both think and do.

3. Shared foundation that grows (reinforced agents)

Knowledge doesn't stay trapped in individual agents. Two motions create compounding:

- **Coalescence:** Best practices from scattered agents consolidate into reusable building blocks
- **Diffusion:** Those building blocks spread to every agent that needs them

Unlike traditional AI that's static after deployment, collaborative intelligence creates a learning layer: agents observe how work gets done, identify successful patterns, surface suggestions to human builders for review. Builders formalize what works. Improvements ripple to all agents that could benefit.

Human in the loop is critical. Agents propose, humans dispose. Builders maintain ownership. The system adapts to your organization's unique way of working — with human judgment. Trust builds through transparency: you see what's suggested, you decide what gets implemented.

Result: Your 50th workflow is easier to build than your 5th because the foundation has grown. Marginal complexity approaches zero as shared knowledge expands. Humans stay in control of what becomes organizational infrastructure.

The competitive landscape

The market hasn't fully shifted to Wave 3. Here's where major players are positioned:

Player	Wave	Focus	Gap
OpenAI	2+	AGI narrative; consumer + large enterprise	No agent-to-agent collaboration or org learning
Microsoft Copilot	2+	Governance-first; tight Office integration	Ecosystem lock-in; limited cross-platform
Google Gemini	2+	Deep Workspace integration	Same ecosystem lock-in problem
Glean	1.5	"Unified search eliminating AI sprawl"	Finds and suggests, but doesn't do or collaborate
Dust	3	Collaborative intelligence layer	Agent↔agent↔human; knowledge compounds

The reframe: Not "which AI tool?" but "which generation of AI?" Not "individual productivity" but "organizational transformation." Not "single-player" but "multiplayer."

SECTION 2

What companies mistake for success

High adoption doesn't mean transformation. Here's what actually matters.

The pattern shows up repeatedly: majority of employees use AI for simple queries (summarization, basic Q&A, translation). When teams attempt advanced use cases requiring sophisticated workflows, basic queries aren't enough. Done poorly, this creates AI slop — noise, inconsistency, lack of quality. Leadership declares "AI success" but transformation isn't happening.

Good use of AI requires practice and craft, not just access to tools.

Look closer: what are people actually using it for? Mostly simple, individual tasks that don't connect to broader workflows. Is organizational capability changing? Can you execute processes faster as a company? Enter new markets? Change your competitive position? Usually: No.

Individual productivity gains are real and necessary. They prove ROI, unlock budget, build momentum. You need them. But they're not sufficient. They're the foundation, not the destination.

Why individual gains don't automatically scale

- Local optimization can create more noise overall (everyone working faster in isolation)
- Improvements don't spread (what works for one person stays with one person)
- Knowledge fragments across tools and teams
- No compounding effect

Productivity gains (necessary)	Transformational change (sufficient)
Individual task acceleration	Organizational capability expansion
Time saved (hours/week)	Work that wasn't possible before
Additive improvement	Multiplicative compounding
Linear gains that plateau	Accelerating gains from network effects
Local optimization	System-level transformation

The question that matters

> *"If we removed AI tomorrow, would our organization's structure and processes be unchanged? Or would removing AI break workflows that now define how we operate?"*

If unchanged: You have individual productivity gains — valuable, but not transformative.
If workflows break: You're building collaborative intelligence — competitive advantage.

The goal: start with individual productivity gains (prove value), build toward transformation (create advantage). It's *and*, not *or*.

SECTION 3

Core capabilities of collaborative intelligence

Collaborative intelligence isn't only about smarter individual tools. It's about creating an organizational nervous system where information flows, humans and AI work as peers, and knowledge compounds.

Multi-specialized-agent collaboration

Specialized agents work with other agents and with humans in flexible patterns.

- @mention agents like you @mention colleagues — natural collaboration patterns
- Sub-agents: agents spawn other agents for multi-step workflows (agent delegates work like a human would)
- Skills: shared capabilities any agent can invoke (like asking a teammate for their expertise)
- Projects: hubs where humans + agents + context + goals work together on specific initiatives

Observability and trust

Trust in collaborative intelligence requires transparency: authorship tracking (what AI created vs. what humans created), human curation markers, audit trails, version history. At first, you trust agents because you trust the people who created them. Over time, you trust agents because you can see their track record.

In practice: from manual to multiplayer

Traditional customer qualification: Rep manually gathers company data, checks criteria, updates CRM, routes to appropriate colleague. Takes 30+ minutes per lead, inconsistent quality.

With collaborative intelligence: Rep @mentions qualification agent in a project. Agent pulls data via enrichment skill, checks criteria via qualification skill, updates CRM via integration tool, routes appropriately via workflow skill. All actions logged with clear authorship markers. Under 2 minutes, consistent, learns from every interaction.

Not one bloated agent trying to do everything. Multiple specialized agents collaborating through shared skills, with full transparency.

This is how Clay, the data enrichment platform valued at $1.5 billion, scaled their GTM Engineering team 4x with Dust — achieving 100% team adoption and saving 58 hours monthly. (Read the full story)

Cross-functional knowledge compounding

This is the exponential unlock.

Two motions create compounding:

Coalescence (knowledge moves inward): Teams build capabilities independently. Best practices emerge from real usage. The system identifies patterns worth formalizing. Excellence consolidates into shared skills.

Diffusion (knowledge spreads outward): Shared skills available to any agent. Update once, improvement spreads everywhere. New agents inherit the full foundation. Marginal cost of each new capability approaches zero.

The horizontal advantage

This is where a horizontal, tool-agnostic approach becomes crucial. Not verticalized — not "AI for sales" separate from "AI for support" separate from "AI for finance." One collaborative intelligence layer across all teams and tools. An abstraction layer that creates intelligent tissue connecting your entire organization, regardless of underlying systems.

Contrast with workflow automation

Traditional automation (like n8n, Gumploop) thinks in arrows and boxes: if this, then that. Rigid, predetermined flows. Breaks when reality doesn't match the flowchart. Adding capability means adding complexity.

Collaborative intelligence thinks in context and skills: agents that reason about situations, flexible and adaptive responses, handles unexpected scenarios. Adding capability means enriching foundation — not adding more spaghetti.

> **The strategic implication:** Companies with this infrastructure build capabilities exponentially faster than competitors starting from scratch each time. Your 50th workflow is easier than your 5th because shared foundation has grown.

Creative Force, an end-to-end content management platform, saw new hires reach full productivity 40% faster after deploying collaborative AI across their organization. (Read the full story)

SECTION 4

Implementation framework

You understand the vision. Now: how do you actually build this?

The framework spans four phases, typically a few weeks to 3 months depending on company size, with ongoing optimization.

Phase	Focus	Timeline
1. Align	Strategy and use case identification	Week 1
2. Foundation	Builder community	Week 1–2
3. Deploy	Build and iterate	Week 2–8
4. Scale	Expand and measure	Week 8+

Finding the right use cases

Don't pick one isolated use case. Pick one use case per team, across at least 3 teams. This reveals what's reusable vs. team-specific, tests whether solutions scale, builds shared foundation from day one, and creates momentum across the organization.

Three lenses for identifying strong use cases:

1. The Anti-To-Do List. Tasks you no longer need to focus on: monitoring KPIs and alerting on changes, tracking competitors and market signals, routine status updates and reporting. Every time you do something annoying, ask "How can I not have to do this again?"

2. Skill bottlenecks. Where work slows when employees hit the limits of their expertise: data analysis when you're not an analyst, trend visualization when you're not a designer, legal review when you're not a lawyer. Where do people wait for experts from other departments?

3. Build new ideas. Unblock thinking and kickstart creation: brainstorm campaign ideas with instant feedback, look for quick insights in raw data, analyze trends without waiting for analysts. Where do people stare at blank pages?

Characteristics of strong first use cases:

- Aligned with a company objective — supports strategic priority, not just "nice to have"

- Understandable by the builder — if you can't understand the problem, you can't test quality
- Currently time-consuming or frustrating — clear before/after, motivated team
- Highly repetitive — compounding is obvious
- Involves 2–3+ people minimum — build trust through collaboration, not isolation

Phase 1: Align on strategy (Week 1)

- Document 3–5 use cases across at least 3 teams (current pain, desired outcome, success metrics, team involved)
- Prioritize using effort/value framework
- Connect to business strategy — which use cases support strategic priorities?
- Build core team: executive sponsor, program lead, 3–5 builders from key departments, admin/IT
- Define success criteria and realistic timeline

Phase 2: Build the builder community (Week 1–2)

Builders are your internal champions who drive implementation. Look for people already using LLMs, with deep understanding of team workflows, a tinkerer mindset, and willingness to share knowledge.

- Create a builder hub: dedicated Slack channel, regular meetups (weekly initially), knowledge sharing space
- Initial training: how to build agents, use skills and tools, test quality, gather feedback and iterate
- Observability best practices: authorship tracking, curation markers

> **Deliverable by end of Phase 2:** 3–5 builders trained and ready to build. Builder community hub established. 3–5 target use cases identified and prioritized across 3+ teams. Workspace setup complete.

Phase 3: Build, deploy, iterate (Week 2–8)

Week 2–4: Builders create 2–3 agents per use case. Keep agents specialized. Make them action-capable from day one. Use shared skills where appropriate. Set up observability.

Week 2–6: Daily usage and rapid iteration. Agents live where work happens (Slack, email, team tools, Projects). Quick fixes daily, weekly reviews of what's working and spreading.

Week 6–8: Measure quantitative metrics (time saved, usage rates, quality). Capture qualitative insights: "What can you do now that you couldn't before?" "Would you be sad if you couldn't use this anymore?"

Phase 4: Scale and optimize (Week 8+)

Extract and formalize shared patterns (Week 8–10): What logic is universal across teams? Create shared skills. Where do workflows connect? Build cross-team agents. What improvements ripple to multiple teams? Formalize and spread.

Expand to adjacent teams (Week 10–12): Pick 2–3 new teams, each building on shared foundation plus team-specific needs. Builder-led scaling — pilot builders mentor new builders.

Ongoing optimization (Week 12+): Monthly builder meetups. Quarterly reviews with leadership on transformation metrics. Regular feedback collection. Iterate on governance as you scale.

> **The compounding effect becomes visible:** Second wave of agents built faster than first. Improvements in one domain lift others automatically. Marginal complexity of new agents decreases. Foundation gets richer without getting more complex.

Long-term vision (3–6 months):

- 20–50+ agents across the organization
- 10–20 shared skills
- Knowledge compounding visible — improvements spread automatically
- New hires onboard to "work includes AI" (not "here's an AI tool to try")
- If you removed AI, workflows would break — it's infrastructure now

SECTION 5

Measuring success

ROI isn't one number. It's a framework tailored to your context.

Before measuring anything, answer: What are your strategic priorities? How would AI supporting those priorities show up? What metrics do your executives already care about?

Framework 1: Internal satisfaction and sentiment

If people don't value it, it's not working — regardless of metrics.

> *"Would you be sad if you couldn't use this anymore?"*

Run quarterly pulse surveys (3–5 questions, keep short). Track sentiment over time. Segment by team and role.

Framework 2: Time and capacity measurement

For each use case, compare before and after: how long does it take today? How often is it performed? New time required? Quality different?

> **Example:** Lead qualification. Before: 30 min/lead × 100 leads/week × 5 AEs = 250 hours/week. After: 5 min/lead = 42 hours/week. Impact: 208 hours/week = 5.2 FTE capacity unlocked.

Framework 3: Organizational capability and strategic impact

Run the "remove AI tomorrow" test. Track what new work is now enabled, how fast work moves between teams, how many agents share common skills, and how quickly improvements spread.

Framework 4: The economic Turing test

Assess AI like you'd assess an employee. "If this were a person, would you hire them? Promote them? Give them more responsibility?" Rate on reliability, speed, quality, learning, collaboration, scope. Compare to human benchmark. Track progression over time.

Your priority	Recommended frameworks
Proving ROI for budget	Time/capacity + Strategic impact
Scaling adoption	Satisfaction + Economic Turing test
Board/exec reporting	Capability/strategic + Quantitative backup
Operational improvement	Efficiency + Quality assessment

Start simple, add depth over time: Week 1–4 basic satisfaction and usage. Week 4–12 time savings and task completion. Month 3–6 new capabilities and strategic impact. Month 6+ knowledge compounding and competitive advantage.

SECTION 6

The path forward

Every company will have AI by 2027. The question is: what kind?

Path 1: Individual productivity AI	Path 2: Collaborative intelligence
Tools that make individuals faster	Systems that make organizations more capable
Vertical solutions, fragmented knowledge	Unified platform, knowledge compounds
Single-player mode	Multiplayer mode
Linear gains that plateau	Exponential gains that accelerate
Competitive parity at best	Sustainable competitive advantage

The gap between these paths is growing. And it compounds.

What to do next

If you're just starting:

- Identify builders across at least 3 teams (who understands workflows + has curiosity?)
- Use the three lenses: Anti-To-Do List, Skill bottlenecks, Build new ideas
- Pick one use case per team, across 3+ teams — build shared foundation from start
- Start building week 1 (don't wait, building is simple)

If you're scaling existing AI:

- Run the "remove AI tomorrow" test: would workflows break?
- Audit what you have: is knowledge compounding or fragmenting?
- Look for patterns: what capabilities should be shared vs. team-specific?
- Build the foundation: extract shared skills, enable cross-team collaboration

If you're transforming your organization:

- Measure beyond individual productivity: what capabilities exist now that didn't before?
- Build the learning layer: how do agents suggest improvements humans can approve?
- Think infrastructure: how does collaborative intelligence become how you operate?
- Track the compounding: is your 50th workflow easier than your 5th?

The companies that will win

They won't be the ones with the smartest AI.

They'll be the ones where AI, humans, and organizational knowledge work as a unified system. Where insights from one team automatically improve another. Where knowledge compounds instead of fragments. Where the organization itself learns and adapts. Where removing AI would break workflows that define how they operate.

That's collaborative intelligence. That's the frontier.

And it's being built now.

dust.tt

The AI operating system for the enterprise.

References

Clay: How Clay is powering 4x team growth with Dust

How Creative Force is using Dust to define the future of work

Why Mirakl chose Dust as its go-to agentic solution

How Alan scaled analytics with Dust